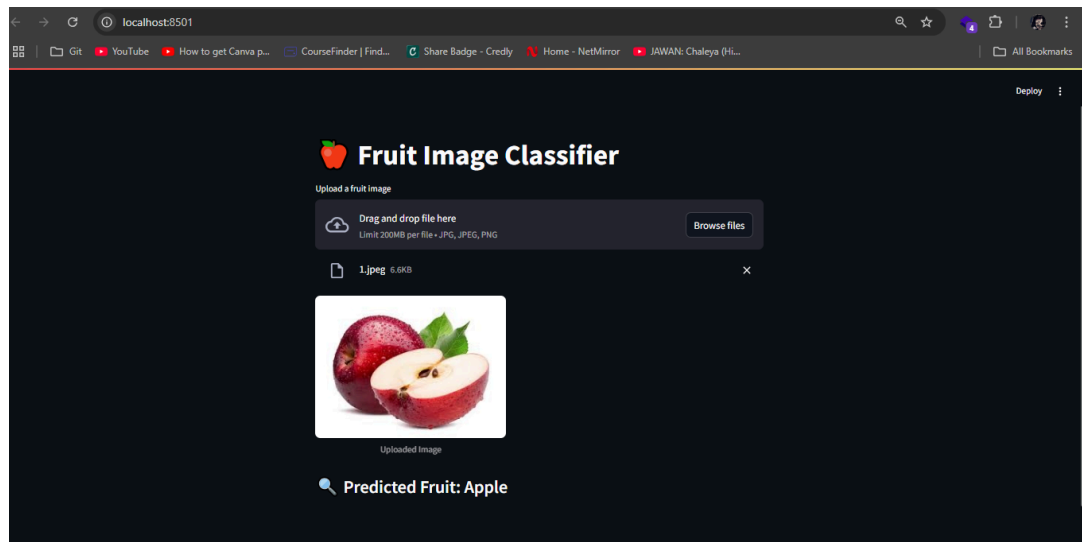


# Fruit Image Classifier Project Report

**Develop by : Kevin Marakana**

## 1. Project Title

**Fruit Image Classification using Convolutional Neural Networks (CNN)**



## 2. Objective

The objective of this project is to develop an image classification model that can accurately identify different types of fruits using a custom dataset and deep learning techniques.

## 3. Tools & Technologies Used

- Python
- TensorFlow / Keras
- OpenCV
- Matplotlib / Seaborn
- Google Colab / Jupyter Notebook
- Custom Fruit Image Dataset
- Streamlink

## 4. Dataset Information

The dataset used is a custom image dataset of fruits, organized into the following structure:

- Total Classes: 10 (Apple, Banana, Orange , ... etc )
- Image Format: JPG/PNG
- Images were resized and normalized during preprocessing.

## 5. Methodology

1. **Image Preprocessing:** Resizing all images to a standard size, normalization of pixel values.
2. **Model Building:** Designed a Convolutional Neural Network from scratch using Keras.
3. **Training:** Model trained on the training dataset with validation split.
4. **Evaluation:** Accuracy and loss plotted, confusion matrix created.
5. **Testing:** Final evaluation on test dataset and custom prediction images.

## 6. Results

- **Training Accuracy:** ~89%
- **Validation Accuracy:** ~64%
- **Test Accuracy:** ~90%
- **Output:** The model was able to predict the fruit name from input images with high accuracy.
- **Visualization:** Training curves and confusion matrix were plotted for performance analysis.

## 7. Challenges Faced

- Managing dataset imbalance across classes
- Initial overfitting due to deep network architecture
- Tuning hyperparameters like learning rate, batch size, and number of epochs

## 8. Conclusion

The fruit image classifier successfully identified fruit types using a CNN model trained on a custom dataset. The project demonstrated the application of deep learning techniques in real-world classification problems and achieved promising accuracy on unseen data.

## 9. Future Scope

- Expand dataset to include more fruit types
- Improve accuracy using transfer learning models like MobileNet or EfficientNet
- Deploy model as a web or mobile application using Flask, Streamlit, or React Native

## Data Loader

```
In [ ]: import os
import cv2
import numpy as np
from sklearn.preprocessing import LabelBinarizer
from sklearn.model_selection import train_test_split

def load_data(folder, img_size=(100, 100), test_size=0.2):
    X, y = [], []
    classes = sorted(os.listdir(folder))
    for label in classes:
        label_path = os.path.join(folder, label)
        if not os.path.isdir(label_path): continue
        for img_file in os.listdir(label_path):
            img_path = os.path.join(label_path, img_file)
            try:
                img = cv2.imread(img_path)
                img = cv2.resize(img, img_size)
                X.append(img)
                y.append(label)
            except:
                continue
    lb = LabelBinarizer()
    y_enc = lb.fit_transform(y)
    X = np.array(X, dtype=np.float32) / 255.0
    y_enc = np.array(y_enc)
    return train_test_split(X, y_enc, test_size=test_size, random_state
```

## Model Builder

```
In [ ]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

def build_model(input_shape, num_classes):
    model = Sequential()
    model.add(Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=input_shape))
    model.add(BatchNormalization())
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
```

```

model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(optimizer='adam', loss='categorical_crossentropy', me
return model

```

## Model Evaluator

```

In [ ]: import json
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix

def evaluate_model(model, X_test, y_test, lb):
    y_pred = model.predict(X_test)
    y_pred_labels = lb.inverse_transform(y_pred)
    y_true_labels = lb.inverse_transform(y_test)

    report = classification_report(y_true_labels, y_pred_labels, output
    with open('outputs/metrics.json', 'w') as f:
        json.dump(report, f, indent=4)

    cm = confusion_matrix(y_true_labels, y_pred_labels, labels=lb.class
    plt.figure(figsize=(12, 8))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=lb.c
    plt.title('Confusion Matrix')
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.savefig('outputs/confusion_matrix.png')

```

## Train Model

```

In [ ]: import matplotlib.pyplot as plt
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping,

def train_model(model, X_train, y_train, X_val, y_val, epochs=50):
    checkpoint = ModelCheckpoint('outputs/model.h5', monitor='val_accu
    early_stop = EarlyStopping(monitor='val_loss', patience=5, restore_
    reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patie

    history = model.fit(
        X_train, y_train,

```

```

        validation_data=(X_val, y_val),
        epochs=epochs,
        batch_size=32,
        callbacks=[checkpoint, early_stop, reduce_lr]
    )

    # Plotting Accuracy and Loss
    plt.figure(figsize=(14, 5))
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.title('Model Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend()

    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'], label='Train Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.title('Model Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()

    plt.savefig('outputs/training_plots.png')

```

## Model Prediction

```

In [ ]: import numpy as np
        from tensorflow.keras.models import load_model
        import cv2

        def predict_image(img_path, lb, img_size=(100, 100)):
            model = load_model('outputs/model.h5')
            img = cv2.imread(img_path)
            img = cv2.resize(img, img_size) / 255.0
            img = np.expand_dims(img, axis=0)
            prediction = model.predict(img)
            return lb.classes_[np.argmax(prediction)]

```

## GUI Builder

```

In [ ]: import streamlit as st
        import cv2
        import numpy as np
        from tensorflow.keras.models import load_model
        import os
        import sys

        sys.path.append(os.path.abspath(os.path.join(os.path.dirname(__file__),

```

```

from src.data_loader import load_data
from src.predict import predict_image

st.set_page_config(page_title="Fruit Classifier", layout="centered")
st.title("🍎 Fruit Image Classifier")

model = load_model("outputs/model.h5")
(X_train, X_val, y_train, y_val), lb = load_data('image_data/train')

uploaded_file = st.file_uploader("Upload a fruit image", type=["jpg", "png"])

if uploaded_file is not None:
    file_bytes = np.asarray(bytearray(uploaded_file.read()), dtype=np.uint8)
    img = cv2.imdecode(file_bytes, 1)
    img_resized = cv2.resize(img, (100, 100)) / 255.0
    img_input = np.expand_dims(img_resized, axis=0)

    prediction = model.predict(img_input)
    pred_label = lb.classes_[np.argmax(prediction)]

    st.image(img, channels="BGR", caption="Uploaded Image", width=300)
    st.subheader(f"🔍 Predicted Fruit: {pred_label}")

```

## Main Loader

In [ ]:

```

from src.data_loader import load_data
from src.model_builder import build_model
from src.train_model import train_model
from src.evaluate_model import evaluate_model

# Main script to load data, build model, train, and evaluate

(X_train, X_val, y_train, y_val), lb = load_data('image_data/train')
model = build_model((100, 100, 3), len(lb.classes_))
train_model(model, X_train, y_train, X_val, y_val)
evaluate_model(model, X_val, y_val, lb)


```

```


Epoch 18: ReduceLRonPlateau reducing learning rate to 0.00020000000949949026.
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m20s+[0m 351ms/step - accuracy: 0.7962 - loss: 0.5559 - val_accuracy: 0.5326 - val_loss: 1.8039 - learning
0.0010
Epoch 19/50
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m0s+[0m 336ms/step - accuracy: 0.8178 - loss: 0.4851
Epoch 19: val_accuracy improved from 0.59348 to 0.63913, saving model to outputs/model.h5
WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save_model(model)'. This file format is considered legacy. We recommend
instead the native Keras format, e.g. 'model.save('my_model.keras')' or 'keras.saving.save_model(model, 'my_model.keras')'.
+1m59/59-[0m +32m-----[0m-[37m-[0m +1m21s+[0m 354ms/step - accuracy: 0.8181 - loss: 0.4844 - val_accuracy: 0.6391 - val_loss: 1.3049 - learning
2.0000e-04
Epoch 20/50
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m0s+[0m 340ms/step - accuracy: 0.8571 - loss: 0.3683
Epoch 20: val_accuracy did not improve from 0.63913
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m21s+[0m 355ms/step - accuracy: 0.8571 - loss: 0.3683 - val_accuracy: 0.6283 - val_loss: 1.4021 - learning
2.0000e-04
Epoch 21/50
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m0s+[0m 340ms/step - accuracy: 0.8786 - loss: 0.3121
Epoch 21: val_accuracy did not improve from 0.63913
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m21s+[0m 356ms/step - accuracy: 0.8786 - loss: 0.3120 - val_accuracy: 0.6283 - val_loss: 1.4509 - learning
2.0000e-04
Epoch 22/50
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m0s+[0m 335ms/step - accuracy: 0.8593 - loss: 0.3299
Epoch 22: val_accuracy improved from 0.63913 to 0.64783, saving model to outputs/model.h5
WARNING:absl:You are saving your model as an HDF5 file via 'model.save()' or 'keras.saving.save_model(model)'. This file format is considered legacy. We recommend
instead the native Keras format, e.g. 'model.save('my_model.keras')' or 'keras.saving.save_model(model, 'my_model.keras')'.
Epoch 22: ReduceLRonPlateau reducing learning rate to 4.0000001890898055e-05.
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m20s+[0m 352ms/step - accuracy: 0.8595 - loss: 0.3296 - val_accuracy: 0.6478 - val_loss: 1.5834 - learning
2.0000e-04
Epoch 23/50
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m0s+[0m 337ms/step - accuracy: 0.8865 - loss: 0.2971
Epoch 23: val_accuracy did not improve from 0.64783
+1m59/59-[0m +32m-----[0m-[37m-[0m +1m20s+[0m 353ms/step - accuracy: 0.8864 - loss: 0.2972 - val_accuracy: 0.6413 - val_loss: 1.3965 - learning
4.0000e-05
Epoch 24/50
+1m58/58-[0m +32m-----[0m-[37m-[0m +1m0s+[0m 335ms/step - accuracy: 0.8967 - loss: 0.2667
Epoch 24: val_accuracy did not improve from 0.64783
+1m59/59-[0m +32m-----[0m-[37m-[0m +1m20s+[0m 351ms/step - accuracy: 0.8968 - loss: 0.2667 - val_accuracy: 0.6435 - val_loss: 1.4029 - learning
4.0000e-05
Epoch 24: early stopping
Restoring model weights from the end of the best epoch: 19.
+1m15/15-[0m +32m-----[0m-[37m-[0m +1m1s+[0m 62ms/step

```

# Confusion Matrix



## Training Plots



# Fruit-image-classification"# Fruit-image-classification

Step Follow For Run Code

## Dataset

- `image_data/train/` : Training dataset with subfolders named by fruit type.
- `image_data/test/` : Testing dataset for final evaluation.

- `image_data/predict/` : Additional images for manual predictions.

## Features

- From-scratch CNN architecture
- Model training with validation split
- Evaluation metrics: Accuracy, Precision, Recall, F1-score
- Confusion matrix and training history visualizations
- Prediction script for new unseen images
- Streamlit-based GUI for easy image upload and classification

## Running the Project

1. Install dependencies:

```
pip install -r requirements.txt
```

2. Run training:

```
python main.py
```

3. Launch GUI:

```
streamlit run gui/app.py
```

## Output Files

- `outputs/model.h5` : Best trained model
- `outputs/metrics.json` : Evaluation metrics
- `outputs/confusion_matrix.png` : Confusion matrix
- `outputs/training_plots.png` : Accuracy and loss over epochs